

Developing AntBot: Visual Navigation based on the insect brain

Robert Mitchell

Master of Informatics
Informatics
School of Informatics
The University of Edinburgh
January 15, 2018

Supervised by
Dr. Barbara Webb

Acknowledgements

I would like to take the opportunity to thank my supervisor, Dr. Barbara Webb, for her guidance, and invaluable insight on the subject matter. My gratitude also extends to Zhaoyu Zhang and Leonard Eberding, two of my predecessors on this project; both have been extremely useful in explaining the existing codebase and operations of the robot where they were not always clear. Finally I should like to thank my parents, for their unwavering support throughout my education; I could not have made it here without them.

Abstract

Declaration

I declare that this dissertation was composed by myself, the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

Robert Mitchell

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Goals	2
1.3	Results	3
2	Background	4
2.1	Optical Flow	4
2.2	Optical flow models for Collision Avoidance	4
2.2.1	Time-to-Collision	4
2.2.2	Filtering	6
2.3	The Mushroom Body for Visual Navigation	6
3	Platform	9
4	Methods	10

List of Figures

- 1 The OF filter model: (Caption from *Stewart et al.* Figure 7): Collision avoidance (CA). (A) The CA filters used in the model. Each covers 105 deg. of azimuth but they are centred at ± 3 deg. (elevation = 0 deg.). (B) Control diagram for collision avoidance. Only the half of the system that triggers rightward saccades is shown for clarity; the other half has an identical configuration. The dark blue box represents the blue wide-field filter in A. The reset operation also applies to the other half of the system, i.e. a saccade in one direction sets both accumulators to 0. Thresh., threshold; ang. vel., angular velocity. 7

- 2 The Mushroom Body circuit: (Caption from *Stewart et al.* Figure 2): Images (see Fig 1) activate the visual projection neurons (vPNs). Each Kenyon cell (KC) receives input from 10 (random) vPNs and exceeds firing threshold only for coincident activation from several vPNs, thus images are encoded as a sparse pattern of KC activation. All KCs converge on a single extrinsic neuron (EN) and if activation coincides with a reward signal, the connection strength is decreased. After training the EN output to previously rewarded (familiar) images is few or no spikes. 8

List of Tables

1 Introduction

Desert ants (*Cataglyphis velox*) have the remarkable ability to navigate through complex natural environments, using only low-resolution visual information and limited computational power. It is well documented that many species of ant, and other hymenoptera are capable of very robust visual navigation; however, it is as yet unclear how the insects perform this seemingly complex task with such little brainpower. In this paper, we will focus on using and extending an existing model for visual navigation in ants using the Mushroom Body circuit, an artificial neural network which emulates the Mushroom Body neuropils in the ant brain. We will also discuss biologically plausible methods of visual Collision Avoidance using Optical Flow. A robot (AntBot) has been constructed [?] to allow us a testing platform on which to implement, and experiment with, the algorithms in the *Ant Navigational Toolkit* [?].

1.1 Motivation

Though we are able to observe, and mimic algorithmically, the visual navigational capabilities of insects, we still do not understand the precise methods by which this process takes place. The model we will look at was proposed by *Ardin et al.* [?], which takes the Mushroom Body (whose function was thought to be primarily for olfactory learning), and shows that this provides a plausible neural model for encoding visual memories.

The MB circuit has been implemented and tested on AntBot by Eberding and Zhang respectively, however the existing MB circuit is fairly simple. It uses binary weightings for the connections between the visual projection neurons and the Kenyon Cells, and a single boolean Extrinsic Neuron denoting image recognition. A modification was made by Zhang, whereby eight ENs were used, one for each of the cardinal directions in the Central Complex model. This will be discussed further in 2.3. The reader should note that the Central Complex (CX) model is primarily used to model the task of Path Integration and will not be discussed further (see [?]).

We would also like to look at methods for collision avoidance (CA) which do not involve specialised sensors such as a LIDAR or SONAR, the luxury of which, ants do not have. Models have been proposed which use

Optical Flow (OF) properties to determine whether or not a collision is imminent. These models have been proposed both in purely robotic contexts [?], and biological ones [?].

1.2 Goals

The project aims for the following experimental scenario to be possible: We want to send the robot on a run through an obstacle course, allowing it to navigate however it chooses through the environment. From here, we want the robot to be able to replicate this route using only visual memories, which it should store on that initial run. Finally, we would like the robot to be able to navigate home following the reverse of this route. It should be noted that this final step is not strictly accurate to the behaviour of the desert ant. *Wehner et al.* [?] demonstrated that the remembered routes have a distinct polarity, so knowledge of a route from nest to food, does not imply that the ant has knowledge of a route from food to nest. In this case, we make the outward and homeward route the same.

The first stage of the project will focus upon obtaining a working collision avoidance system as a pre-requisite to gathering the route information. This CA system should be based on visual information readily available to AntBot with no additional/specialist sensors. For this paper, we assume that CA is a low-level reactionary behaviour, in that, we do not use any further processing of the detected motion; we react based on the immediate stimulus of the flow field. We will look at two different optical flow techniques used to build CA systems. We will also discuss the effects of using different types of flow field, how the different flow techniques behave in the same situation, and different methods of response.

We then move to the Mushroom Body circuit, first establishing a baseline performance measure for the visual navigation task by using the original *basic* model from [?] with binary weightings and a single Extrinsic Neuron. A scanning behaviour will be used for this baseline; ants have demonstrated use of scanning in visual navigation but it is generally accepted that this is not the primary method they use to determine a direction after having recognised a scene, rather, this scanning behaviour only occurs in certain scenarios (e.g. when the ant becomes lost) [?].

Finally, we will report the results of the experiments performed at different stages during, and post development; we will compare these to relevant results from previous iterations of this project. We will end with a conclusion of our findings and contributions to the project, as well as discussing technical limitations and potential for future developments.

1.3 Results

This work is based on work done previously by Leonard Eberding, Luca Scimeca, and Zhaoyu Zhang [?, ?, ?].

Significant contributions:

1. An optical flow based system for Collision Avoidance
2. Results indicating the impracticality of an expansion based system for Collision Avoidance
3. *[FUTURE]* Successful replication of a route through a cluttered environment using Visual Navigation
4. *[FUTURE]* Comparison of different Visual Navigation models in a set navigational task

Note: results marked [FUTURE] are results I would like to have achieved by the culmination of the project and the writing of the final dissertation. They are not current. To be continued...

2 Background

2.1 Optical Flow

Image flow is defined as being the 3D velocity vector of an object, projected onto a 2D image plane[?]. Optical flow is an approximation of this, working from a series of images to compute the projected velocity vector for a pixel. A single *flow vector* shows the displacement of a single pixel from one image to the next. A set of these flow vectors creates a *flow field*, a series of vectors which describe the motion in the complete image.

Broadly, there are two types of flow field: dense and sparse (also known as differential and feature-based respectively[?]). A *dense* flow field tracks the motion of every pixel in the image. A *sparse* flow field tracks the motion of a subset of pixels in the image. The sparse field may track a uniform subset of pixels, such as a grid, or a set of important points in the image, such as the prominent features in the image (object corners). See 4 for implementation.

2.2 Optical flow models for Collision Avoidance

Collision avoidance is an important component in navigation. Ants do not have dedicated sensory systems or the ability to create a visual 3D map of their environment using stereoscopic vision or motion parallax. It has been demonstrated that optical flow is used by honeybees in performing visual navigation [?], so it is not unreasonable to think that bees and other hymenoptera may use this information for other purposes. Indeed, optical flow has been shown to be a viable model for collision avoidance in *Drosophila* [?]. We visit two models in this paper which we will term the *time-to-collision* model and the *filtering* model.

2.2.1 Time-to-Collision

This model, proposed by *Low and Wyeth*[?], aims to replicate the CA capabilities of larger animals such as birds and humans. The model relies on accurately computing the *time-to-collision* or *time-to-contact* (TTC), which is the computed time until the robot collides with an obstacle. The TTC is computed as follows:

$$TTC = \frac{d}{v} \quad (1)$$

Where d is the distance from a point object on a collision course with the robot and v is constant speed at which the robot and object are closing. *Low and Wyeth* then alter this equation by taking the direction of the velocity vector v and the direction to a point on the object and computing the angle between these two vectors ϕ . Their TTC equation then becomes:

$$TTC = \frac{\cos\phi \times \sin\phi}{\phi} \quad (2)$$

However, this method of computation relies on the assumption that the robot only ever moves in the forward camera direction which is not appropriate for AntBot (See section 3). We must look for an alternate way to compute TTC. Fortunately, such information is easily computable from the *focus of expansion* (FOE), the point from which all flow vectors originate.

A general method for computing the FOE is given by [?]:

$$FOE = (A^T A)^{-1} A^T \mathbf{b} \quad (3)$$

$$A = \begin{bmatrix} a_{00} & a_{01} \\ \dots & \dots \\ a_{n0} & a_{n1} \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_0 \\ \dots \\ b_n \end{bmatrix}$$

Where, for each pixel $p_i = (x, y)$, the associated flow vector is given by $\mathbf{v} = (u, v)$. We then set $a_{i0} = u$, $a_{i1} = v$ and finally $b_i = xv - yu$. The TTC can then be computed as:

$$TTC = \frac{d}{v} = \frac{y}{\frac{\partial y}{\partial t}} \quad (4)$$

Where y is the vertical distance of some point $p = (x, y, z)$ from the FOE, and $\frac{\partial y}{\partial t}$ is the velocity of translational motion of y . A full derivation is given by *O'Donovan* from whom we have adapted this equation.

Finally, we can simplify this to:

$$TTC = \frac{\Delta_i}{|\vec{V}_t|} \dots \quad (5)$$

Where Δ_i is the distance of a point $p_i = (x, y)$ from the FOE, and $|\vec{V}_t|$ is the translational velocity of the camerap computed from optical flow[?]. This time-to-contact is then appropriately thresholded and a reaction is generated based on the position of the focus of expansion.

2.2.2 Filtering

The filtering method asks the following question: Given my current motion, what visual changes do I expect to see? Much of the following explanation was provided by [?]. A model proposed by *Stewart et al.* for CA in simulated fruit flies takes advantage of the fact that expanding patterns will trigger an avoidance manouver away from the focus of expansion[?]. Their model uses two offset flow filters (the *expected* flow). Each filter is constructed as a frontally centered expansion pattern with the same spatial extent to either side of the expansion pole (the central vertical axis of the pattern). These filters are then offset by $+3^\circ$ for the right and -3° for the left. The left and right filters feed into leaky accumulators, and if the accumulator exceeds a given threshold, a saccade in the opposite direction is triggered. The model from [?] is shown in Figure 1. We discuss a slightly modified version of this model in section 4.

2.3 The Mushroom Body for Visual Navigation

The Mushroom Body neuropils are structures present in the brains of insects though they are largest in the brains of hymenoptera. They are known to play a critical role in olfactory learning, and have been thought to play a role in visual memory in hymenoptera since 1982 at the latest[?]. In 2010 a Mushroom Body circuit was proposed by *Ardin et al.* to allow emulation of the structure in a simulated desert ant [?]. The simulated ant’s view is taken from 1cm above the ground and has a field of view of 296° azimuth by 76° elevation. A ratio of $4^\circ/\text{pixel}$ is used to give a 19×74 pixel image. This image is then downsampled to 10×36 pixels to give a realistic resolution for ant vision. A 1×360 vector is used for further processing.




Figure 1: The OF filter model: (Caption from *Stewart et al.* Figure 7): Collision avoidance (CA). (A) The CA filters used in the model. Each covers 105 deg. of azimuth but they are centred at ± 3 deg. (elevation = 0 deg.). (B) Control diagram for collision avoidance. Only the half of the system that triggers rightward saccades is shown for clarity; the other half has an identical configuration. The dark blue box represents the blue wide-field filter in A. The reset operation also applies to the other half of the system, i.e. a saccade in one direction sets both accumulators to 0. Thresh., threshold; ang. vel., angular velocity.

The generalised MB circuit is a three layer neural network: The first layer consists of a set of visual Projection Neurons (vPNs), these connect to the second layer of standard artificial neurons referred to as Kenyon Cells (KCs), and finally these KCs connect to a set of Extrinsic Neurons (ENs).

The model by *Ardin et al.* (shown in Figure 2) consisted of 350 vPNs (one for each pixel in the downsampled image). In the second layer we have 20,000 KCs each of which receives input from 10 randomly selected vPNs; each KC requires coincident input from multiple vPNs to fire. Every KC is then connected to a single EN which sums the number of KCs which are activated by the input image. The network is trained by providing a reward signal at regular intervals. If KC activation coincides with a reward signal, the connection strength to the EN is greatly

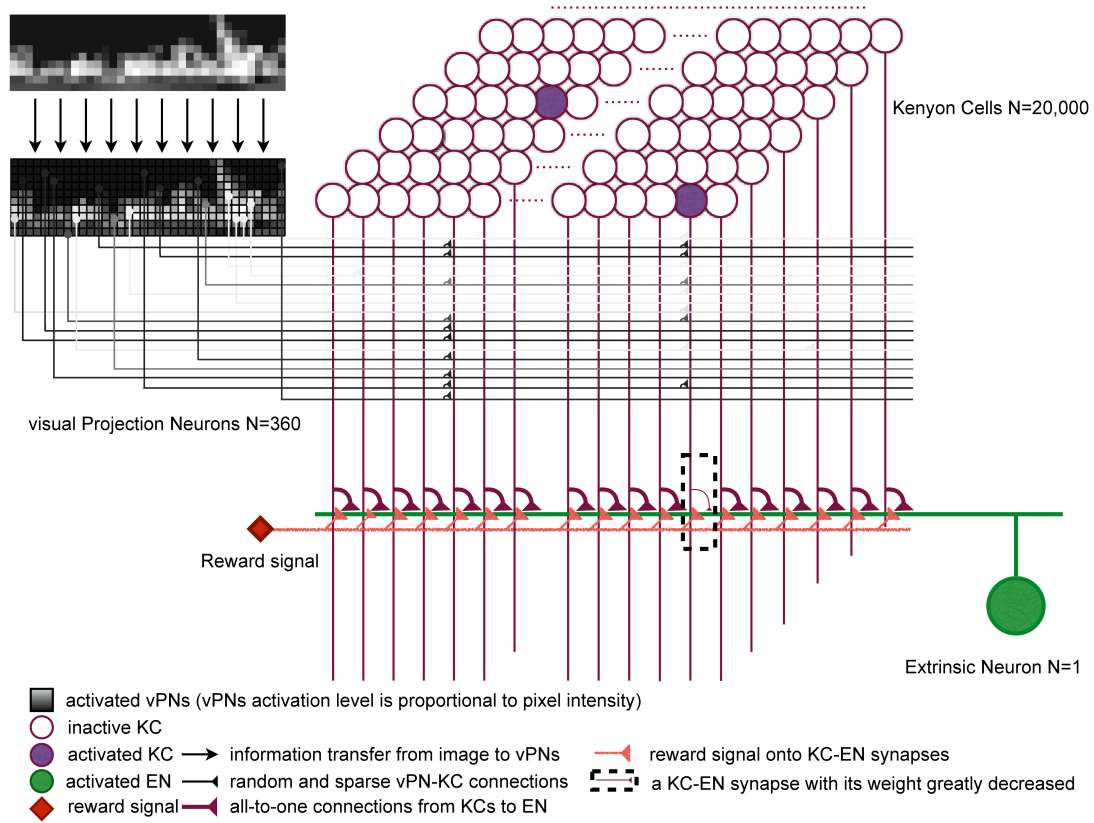


Figure 2: The Mushroom Body circuit: (Caption from *Stewart et al.* Figure 2): Images (see Fig 1) activate the visual projection neurons (vPNs). Each Kenyon cell (KC) receives input from 10 (random) vPNs and exceeds firing threshold only for coincident activation from several vPNs, thus images are encoded as a sparse pattern of KC activation. All KCs converge on a single extrinsic neuron (EN) and if activation coincides with a reward signal, the connection strength is decreased. After training the EN output to previously rewarded (familiar) images is few or no spikes.

reduced. The single EN simply gives a familiarity measure for the image seen. The agent decides on its next action by scanning to find the direction of greatest familiarity.

In 2017, a model was proposed by *Zhang*

3 Platform

a

4 Methods

a